

CSC413: Homework 2

Tianyu Du (1003801647)

2020/02/08 at 22:21:00

1 Optimization

1.1 Stochastic Gradient Descent (SGD)

1.1.1 Minimum Norm Solution

Answer. Recall from the previous homework that the solution found by gradient descent in the over-parameterized situation was

$$\mathbf{w}^* = X^T (X X^T)^{-1} \mathbf{t} \quad (1.1)$$

Moreover, the solution \mathbf{w}^* reached by gradient descent was in the row space of X (i.e., the span of rows of X), and \mathbf{w}^* is a zero loss solution.

For one single updating using \mathbf{x}_i in the t^{th} iteration in the SGD process, the gradient is

$$\nabla_{\mathbf{w}_t} \mathcal{L}_i = \frac{\partial}{\partial \mathbf{w}_t} \|\mathbf{w}_t^T \mathbf{x}_i - t_i\|_2^2 \quad (1.2)$$

$$= 2 \underbrace{(\mathbf{w}_t^T \mathbf{x}_i - t_i)}_{\in \mathbb{R}} \mathbf{x}_i \quad (1.3)$$

$$\implies -\eta \nabla_{\mathbf{w}_t} \mathcal{L}_i \in \text{Row}(X) \quad (1.4)$$

Provided that the starting point $\mathbf{w}_0 = \mathbf{0} \in \text{Row}(X)$, for every t , $\mathbf{w}_t \in \text{Row}(X)$ by an inductive argument. Because $\mathbf{w}_t \in \text{Row}(X)$, let

$$\mathbf{w}_t = X^T \mathbf{r}_t \quad \mathbf{r}_t \in \mathbb{R}^n \quad (1.5)$$

$$\hat{\mathbf{w}} = X^T \hat{\mathbf{r}} \quad \hat{\mathbf{r}} \in \mathbb{R}^n \quad (1.6)$$

Note that, if \mathbf{x}_i is chosen for the t^{th} iteration of SGD, then only the i^{th} components of \mathbf{r}_t and \mathbf{r}_{t+1} are different, and all other components are the same.

Suppose the solution $\hat{\mathbf{w}}$ reached by SGD is a zero loss solution, that is, $X \hat{\mathbf{w}} = \mathbf{t}$. Then

$$X X^T \hat{\mathbf{r}} = \mathbf{t} \quad (1.7)$$

Because $\hat{\mathbf{w}}$ is a global minimum and SGD converges to this point, then it must be the case that

$$\nabla_{\hat{\mathbf{w}}} \mathcal{L}_i(\mathbf{x}_i, \hat{\mathbf{w}}) = 0 \quad \forall i \in \{1, \dots, n\} \quad (1.8)$$

That is, one additional iteration of SGD does not improve performance no matter which sample is chosen. This is the same as

$$\nabla_{r_i} \|XX^T \hat{\mathbf{r}} - \mathbf{t}\|_2^2 = 0 \quad \forall i \in \{1, \dots, n\} \quad (1.9)$$

$$\iff \nabla_{\mathbf{r}} \|XX^T \hat{\mathbf{r}} - \mathbf{t}\|_2^2 = 0 \quad (1.10)$$

$$\implies (XX^T \hat{\mathbf{r}} - \mathbf{t})^T XX^T = 0 \quad (1.11)$$

$$\implies XX^T(XX^T \hat{\mathbf{r}} - \mathbf{t}) = 0^T \quad (1.12)$$

$$\implies XX^T XX^T \hat{\mathbf{r}} = XX^T \mathbf{t} \quad (1.13)$$

$$\implies \hat{\mathbf{r}} = (XX^T)^{-1} \mathbf{t} \quad (1.14)$$

$$\implies \hat{\mathbf{w}} = X^T \hat{\mathbf{r}} = X^T (XX^T)^{-1} \mathbf{t} \quad (1.15)$$

$$= \mathbf{w}^* \quad (1.16)$$

■

1.1.2 Mini-batch SGD (Optional)

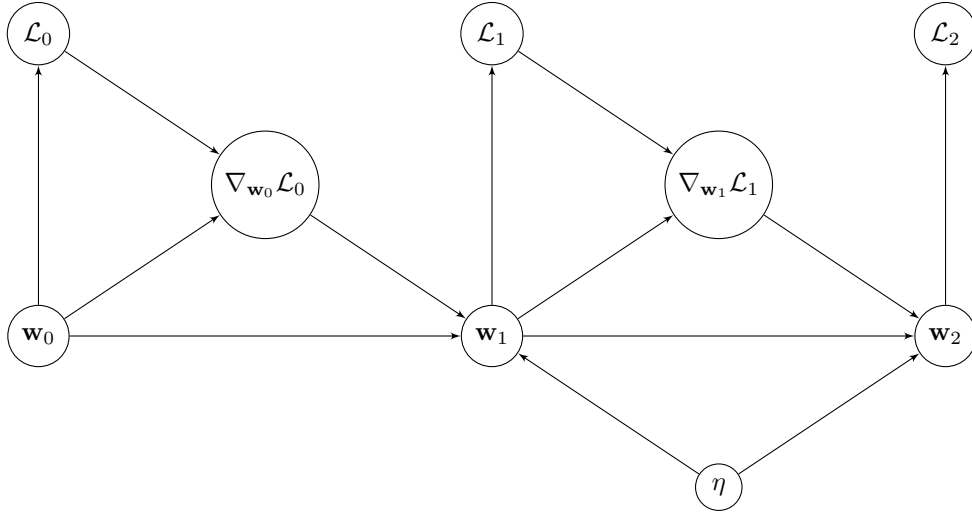
Answer.

■

2 Gradient-Based Hyper-Parameter Optimization

2.1 Computation Graph of Learning Rates

2.1.1



2.1.2

Answer.

(Forward-Propagation) For each iteration t , the prediction $X\hat{\mathbf{w}} \in \mathbb{R}^n$ and label \mathbf{t} take memory of size $2t$, the loss takes 1 unit of memory. Therefore, the overall memory complexity is $\mathcal{O}(d)$.

(**Back-Propagation**) $\nabla_{\eta} \mathcal{L}_t$ needs to be computed using chain rule through \mathbf{w}_{τ} for every $\tau \leq t$. Note that for every τ , $\nabla_{\mathbf{w}_{\tau}} \mathcal{L}_{\tau} \in \mathbb{R}^d$ and $\frac{d}{d\eta} \mathbf{w}_{\tau} \in \mathbb{R}^d$. The memory complexity is $\mathcal{O}(dt)$. ■

2.1.3

Answer. The memory complexity analysis in the previous question suggests the required memory of tracing gradient with respect to η grows linearly in the number of iterations, the memory cost can be prohibitive. ■

2.2 Learning Learning Rates

2.2.1

Answer.

$$\mathbf{w}_1 = \mathbf{w}_0 - \eta \nabla_{\mathbf{w}_0} \mathcal{L}_0 \quad (2.1)$$

$$= \mathbf{w}_0 - \frac{2\eta}{n} X^T (X \mathbf{w}_0 - \mathbf{t}) \quad (2.2)$$

$$\mathcal{L}_1 = \frac{1}{n} \left\| X \left[\mathbf{w}_0 - \frac{2\eta}{n} X^T (X \mathbf{w}_0 - \mathbf{t}) \right] - \mathbf{t} \right\|_2^2 \quad (2.3)$$

$$= \frac{1}{n} \left\| X \left[\mathbf{w}_0 - \frac{2\eta}{n} X^T \mathbf{a} \right] - \mathbf{t} \right\|_2^2 \quad (2.4)$$

$$= \frac{1}{n} \left\| X \mathbf{w}_0 - \frac{2\eta}{n} X X^T \mathbf{a} - \mathbf{t} \right\|_2^2 \quad (2.5)$$

$$= \frac{1}{n} \left\| \mathbf{a} - \frac{2\eta}{n} X X^T \mathbf{a} \right\|_2^2 \quad (2.6)$$

■

2.2.2

Answer. Note that the an arbitrary norm is convex: let $\lambda \in [0, 1]$

$$\|\lambda x + (1 - \lambda)y\|_p \leq \|\lambda x\|_p + \|(1 - \lambda)y\|_p \quad (\text{triangle inequality}) \quad (2.7)$$

$$= |\lambda| \|x\|_p + |1 - \lambda| \|y\|_p \quad (2.8)$$

$$= \lambda \|x\|_p + (1 - \lambda) \|y\|_p \quad (2.9)$$

And $f(x) = x^2$ is strictly convex as well. Moreover, $\mathbf{a} - \frac{2\eta}{n} X X^T \mathbf{a}$ is linear in η . $\mathcal{L}_1(\eta)$ is a composite of linear, convex and strictly convex functions. Therefore, $\mathcal{L}_1(\eta)$ is strictly convex with respect to η . ■

2.2.3

Answer. The derivative of \mathcal{L}_1 w.r.t. η is

$$\nabla_{\eta} \mathcal{L}_1 = \nabla_{\eta} \frac{1}{n} \left\| \mathbf{a} - \frac{2\eta}{n} XX^T \mathbf{a} \right\|_2^2 \quad (2.10)$$

$$= \frac{2}{n} \left(\mathbf{a} - \frac{2\eta}{n} XX^T \mathbf{a} \right)^T XX^T \mathbf{a} \quad (2.11)$$

$$= \frac{2}{n} \left(\mathbf{a}^T - \frac{2\eta}{n} \mathbf{a}^T XX^T \right) XX^T \mathbf{a} \quad (2.12)$$

$$= \frac{2}{n} \left(\mathbf{a}^T XX^T \mathbf{a} - \frac{2\eta}{n} \mathbf{a}^T XX^T XX^T \mathbf{a} \right) \quad (2.13)$$

Set derivative to zero and solve for the optimal learning rate η^* :

$$\nabla_{\eta} \mathcal{L}_1 = 0 \quad (2.14)$$

$$\Rightarrow \frac{2}{n} \left(\mathbf{a}^T XX^T \mathbf{a} - \frac{2\eta}{n} \mathbf{a}^T XX^T XX^T \mathbf{a} \right) = 0 \quad (2.15)$$

$$\Rightarrow \|X^T \mathbf{a}\|_2^2 = \frac{2\eta^*}{n} \|XX^T \mathbf{a}\|_2^2 \quad (2.16)$$

$$\Rightarrow \eta^* = \frac{n \|X^T \mathbf{a}\|_2^2}{2 \|XX^T \mathbf{a}\|_2^2} \quad (2.17)$$

■

3 Convolutional Neural Networks

3.1 Convolutional Filters

Answer.

$$\mathbf{I} * \mathbf{J} = \begin{bmatrix} -1 & 2 & 2 & -2 & 0 \\ -2 & 1 & 0 & 2 & -1 \\ 3 & 0 & 0 & 1 & -1 \\ -2 & 2 & 0 & 2 & -1 \\ 0 & -2 & 3 & -2 & 0 \end{bmatrix} \quad (3.1)$$

This convolutional filter detect edges. ■

3.2 Size of ConvNets

Answer. The table below presents numbers of parameters in all layers.

Layer	Dim. In	Dim. Out	Num. Weights	Num. Bias	Total Params.
Conv3-64	112*112*3	112*112*64	3*3*3*64	64	1,792
Max Pool	112*112*64	56*56*64	0	0	0
Conv3-128	56*56*64	56*56*128	3*3*64*128	128	73,856
Max Pool	56*56*128	28*28*128	0	0	0
Conv3-256	28*28*128	28*28*256	3*3*128*256	256	295,168
Conv3-256	28*28*256	28*28*256	3*3*256*256	256	590,080
Max Pool	28*28*256	14*14*256	0	0	0
FC-1024	14*14*256	1024	14*14*256*1024	1024	51,381,248
FC-100	1024	100	1024*100	100	102,500
Softmax	100	100	0	0	0
Total					52,444,644

■