

CSC413: Homework 3

Tianyu Du (1003801647)

2020/02/29 at 01:40:45

1 Weight Decay

1.1 Under-parameterized Model [0pt]

Solution. Given

$$\mathcal{J}(\hat{\mathbf{w}}) = \frac{1}{2n} \|X\hat{\mathbf{w}} - \mathbf{t}\|_2^2 + \frac{\lambda}{2} \hat{\mathbf{w}}^T \hat{\mathbf{w}} \quad (1.1)$$

$$\mathcal{J}(\hat{\mathbf{w}}) = \frac{1}{2n} \|X\hat{\mathbf{w}} - \mathbf{t}\|_2^2 + \frac{\lambda}{2} \|\hat{\mathbf{w}}\|_2^2 \quad (1.2)$$

The gradient descent converges when $\frac{d}{d\hat{\mathbf{w}}} \mathcal{J}(\hat{\mathbf{w}}) = 0$. Altogether with the fact that $\frac{d}{d\mathbf{x}} \|\mathbf{x}\|_2^2 = 2\mathbf{x}^T$, the training converges if and only if:

$$\frac{d}{d\hat{\mathbf{w}}} \mathcal{J}(\hat{\mathbf{w}}) = \frac{1}{n} (X\hat{\mathbf{w}} - \mathbf{t})^T X + \lambda \hat{\mathbf{w}}^T \quad (1.3)$$

$$= \frac{1}{n} \hat{\mathbf{w}}^T X^T X - \frac{1}{n} \mathbf{t}^T X + \lambda \hat{\mathbf{w}}^T = 0 \quad (\dagger) \quad (1.4)$$

Note that when $d \leq n$, $\text{rank}(X) = d$ implies $X^T X$ is invertible. Suppose $X^T X + n\lambda I$ is invertible as well. Therefore,

$$(\dagger) \implies \left(\hat{\mathbf{w}}^T X^T X + n\lambda \hat{\mathbf{w}}^T \right) = \mathbf{t}^T X \quad (1.5)$$

$$\implies \hat{\mathbf{w}}^T (X^T X + n\lambda I) = \mathbf{t}^T X \quad (1.6)$$

$$\implies \hat{\mathbf{w}}^T = \mathbf{t}^T X (X^T X + n\lambda I)^{-1} \quad (1.7)$$

$$\implies \hat{\mathbf{w}} = (X^T X + n\lambda I)^{-1} X^T \mathbf{t} \quad (1.8)$$

■

1.2 Over-parameterized Model

1.2.1 Warmup: Visualizing Weight Decay [1pt]

Solution. Given

$$\mathbf{x}_1 = (2, 1) \text{ and } t_1 = 2. \quad (1.9)$$

From previous homework, the solution of gradient descent without regularization is

$$\mathbf{w}^* = \left(\frac{4}{5}, \frac{2}{5} \right) \quad (1.10)$$

From the previous part, the gradient descent converges if and only if

$$\frac{d}{d\hat{\mathbf{w}}} \mathcal{J}(\hat{\mathbf{w}}) = \frac{1}{n} (X\hat{\mathbf{w}} - \mathbf{t})^T X + \lambda \hat{\mathbf{w}}^T \quad (1.11)$$

$$= ((2, 1) \cdot (w_1, w_2) - 2)(2, 1) + \lambda(w_1, w_2) \quad (1.12)$$

$$= (2w_1 + w_2 - 2)(2, 1) + (\lambda w_1, \lambda w_2) = 0 \quad (1.13)$$

$$(1.14)$$

Therefore, the solution to gradient descent with weight decay is:

$$\begin{cases} 4w_1 + 2w_2 - 4 + \lambda w_1 &= 0 \\ 2w_1 + w_2 - 2 + \lambda w_2 &= 0 \end{cases} \quad (1.15)$$

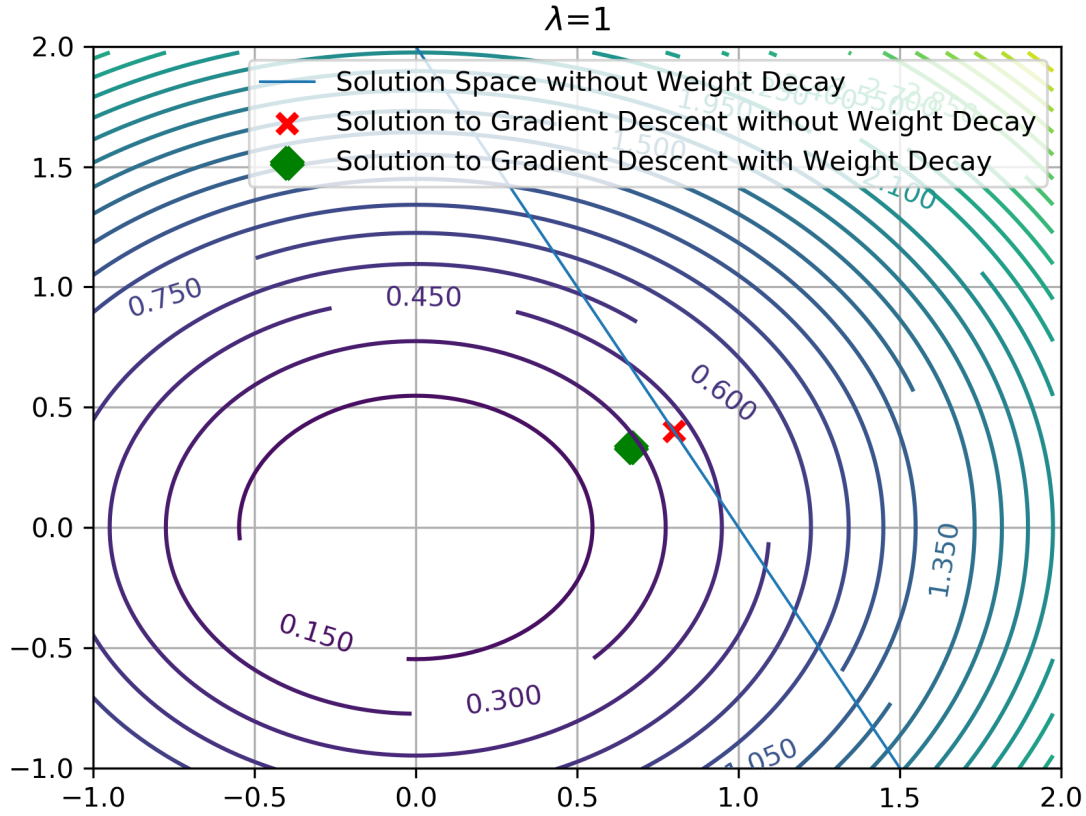
$$\implies \begin{cases} w_1^* &= \frac{4}{\lambda+5} \\ w_2^* &= \frac{2}{\lambda+5} \end{cases} \quad (1.16)$$

The solution space is parameterized by λ as following

$$\mathcal{S} := \left\{ \frac{4}{\lambda+5}, \frac{2}{\lambda+5} : \lambda \in \mathbb{R}_+ \right\} \quad (1.17)$$

which is a singleton uniquely determined by λ . Note that the following visualization assumes $\lambda = 1$.

Figure 1.1: Visualization



■

1.2.2 Gradient Descent and Weight Decay [0pt]

Solution. The solution to gradient descent with weight decay has been derived in the previous section:

$$\mathbf{w}_{\text{weight decay}}^* = \left(\frac{4}{\lambda + 5}, \frac{2}{\lambda + 5} \right) \quad (1.18)$$

■

1.3 Adaptive optimizer and Weight Decay [1pt]

2 Ensembles and Bias-variance Decomposition

2.1 Weight Average or Prediction Average?

2.1.1 [1pt]

Solution. Without loss of generality, assume the bias is zero. This is equivalent to inserting a column of ones to the X , so that $X \in \mathbb{R}^{n \times (d+1)}$, and we can ignore the bias. ■

3 Generalization and Dropout

3.1 Regression Coefficients

3.1.1 Regression from X_1 [0pt]

Solution. ■

3.1.2 Regression from X_2 [1pt]

Solution. Since we are using X_2 only, equivalently, we can set the weight of X_1 to zero:

$$\mathcal{J} = \mathbb{E}_{(x_2, y) \sim (X_2, Y)} [(y - \hat{y})^2] \quad (3.1)$$

$$= \mathbb{E}_{(x_2, y) \sim (X_2, Y)} [(y - w_2 x_2)^2] \quad (3.2)$$

$$= \mathbb{E}_{(x_2, y) \sim (X_2, Y)} [(y - w_2(y + \text{Gaussian}(0, 1)))^2] \quad (3.3)$$

$$= \mathbb{E}_{(x_2, y) \sim (X_2, Y)} [((1 - w_2)y - w_2 \text{Gaussian}(0, 1))^2] \quad (3.4)$$

$$= \mathbb{E}_{(x_2, y) \sim (X_2, Y)} [((1 - w_2)y)^2] + w_2^2 \mathbb{E}_{(x_2, y) \sim (X_2, Y)} [\text{Gaussian}(0, 1)^2] \quad (3.5)$$

$$= (1 - w_2)^2 \mathbb{E}_{y \sim Y} [y^2] + w_2^2 \quad (3.6)$$

Taking the gradient and solve the first order condition:

$$\nabla_{w_2} (1 - w_2)^2 \mathbb{E}_{y \sim Y} [y^2] + w_2^2 = 0 \quad (3.7)$$

$$\implies -2(1 - w_2) \mathbb{E}_{y \sim Y} [y^2] + 2w_2 = 0 \quad (3.8)$$

$$\implies \mathbb{E}_{y \sim Y} [y^2] - w_2 \mathbb{E}_{y \sim Y} [y^2] - w_2 = 0 \quad (3.9)$$

$$\implies \mathbb{E}_{y \sim Y} [y^2] + w_2(1 - \mathbb{E}_{y \sim Y} [y^2]) = 0 \quad (3.10)$$

$$\implies w_2 = \frac{\mathbb{E}_{y \sim Y} [y^2]}{\mathbb{E}_{y \sim Y} [y^2] + 1} \quad (3.11)$$

The expectation of y^2 is

$$\mathbb{E}_{y \sim Y} [y^2] = \mathbb{E}_{x_1 \sim X_1} (x_1 + \text{Gaussian}(0, \sigma^2))^2 \quad (3.12)$$

$$= 2\sigma^2 \quad (3.13)$$

$$\implies w_2 = \frac{2\sigma^2}{2\sigma^2 + 1} \quad (3.14)$$

■

3.1.3 Regression from (X_1, X_2) [1pt]

4 Hard-Coding Recurrent Neural Networks [1pt]

Solution. Let $\sigma = \frac{1}{1 + \exp(-z)}$, and $\mathbf{x}_t = (x_1^t, x_2^t)$ denotes the input feature at time t . Consider the following recurrent network:

$$\hat{y}_t = \sigma(\mathbf{w}_{hy} \mathbf{h}_t + b_y) \quad (4.1)$$

$$\mathbf{h}_t = \sigma(\mathbf{w}_{hh} \mathbf{h}_{t-1} + \mathbf{w}_{xh} \mathbf{x}_t + b_h) \quad (4.2)$$

with the following parameters:

$$\mathbf{w}_{xh} = \begin{pmatrix} a & b \\ c & d \end{pmatrix} \quad (4.3)$$

$$\mathbf{w}_{hh} = \begin{pmatrix} a & b \\ c & d \end{pmatrix} \quad (4.4)$$

$$b_h = 0 \quad (4.5)$$

$$\mathbf{w}_{hy} = \begin{pmatrix} a & b \\ c & d \end{pmatrix} \quad (4.6)$$

$$b_y = 0 \quad (4.7)$$

Justification:

■