

# CSC413: Programming Assignment 4

Tianyu Du (1003801647)

2020/03/26 at 15:57:49

## 1 Part 1: Deep Convolutional GAN (DCGAN) [4pt]

### 1.1 Generator

**Implementation** Please note that I modified the forward method a little bit as well.

```
1 class DCGenerator(nn.Module):
2     def __init__(self, noise_size, conv_dim, spectral_norm=False):
3         super(DCGenerator, self).__init__()
4
5         self.conv_dim = conv_dim # 32
6
7         #####
8         ## FILL THIS IN: CREATE ARCHITECTURE ##
9         #####
10
11        self.linear_bn = nn.Linear(100, self.conv_dim*4*4*4) # (100, 2048)
12        self.upconv1 = upconv(self.conv_dim*4, self.conv_dim*2, 5, stride=2, padding=2,
13                               batch_norm=True, spectral_norm=spectral_norm)
14        self.upconv2 = upconv(self.conv_dim*2, self.conv_dim, 5, stride=2, padding=2,
15                               batch_norm=True, spectral_norm=spectral_norm)
16        self.upconv3 = upconv(self.conv_dim, 3, 5, stride=2, padding=2, batch_norm=True,
17                               spectral_norm=spectral_norm)
18
19    def forward(self, z):
20        """Generates an image given a sample of random noise.
21
22        Input
23        ----
24        z: BS x noise_size x 1 x 1 --> BSx100x1x1 (during training)
25
26        Output
27        -----
28        out: BS x channels x image_width x image_height --> BSx3x32x32 (during
29        training)
30        """
31        batch_size = z.size(0)
32        # Extra reshape
33        z = z.view(batch_size, -1)
34        out = F.relu(self.linear_bn(z)).view(-1, self.conv_dim*4, 4, 4) # BS x 128 x 4 x
35
36        4
37
38        out = F.relu(self.upconv1(out)) # BS x 64 x 8 x 8
39        out = F.relu(self.upconv2(out)) # BS x 32 x 16 x 16
40        out = F.tanh(self.upconv3(out)) # BS x 3 x 32 x 32
41
42        out_size = out.size()
```

```

36         if out_size != torch.Size([batch_size, 3, 32, 32]):
37             raise ValueError("expect {} x 3 x 32 x 32, but get {}".format(batch_size,
out_size))
38         return out
39
40

```

## 1.2 Training Loop

### Implementation

```

1 def gan_training_loop(dataloader, test_dataloader, opts):
2     """Runs the training loop.
3         * Saves checkpoint every opts.checkpoint_every iterations
4         * Saves generated samples every opts.sample_every iterations
5     """
6
7     # Create generators and discriminators
8     G, D = create_model(opts)
9
10    g_params = G.parameters() # Get generator parameters
11    d_params = D.parameters() # Get discriminator parameters
12
13    # Create optimizers for the generators and discriminators
14    g_optimizer = optim.Adam(g_params, opts.lr, [opts.beta1, opts.beta2])
15    d_optimizer = optim.Adam(d_params, opts.lr * 2., [opts.beta1, opts.beta2])
16
17    train_iter = iter(dataloader)
18
19    test_iter = iter(test_dataloader)
20
21    # Get some fixed data from domains X and Y for sampling. These are images that are held
22    # constant throughout training, that allow us to inspect the model's performance.
23    fixed_noise = sample_noise(100, opts.noise_size) # 100 x noise_size x 1 x 1
24
25    iter_per_epoch = len(train_iter)
26    total_train_iters = opts.train_iters
27
28    losses = {"iteration": [], "D_fake_loss": [], "D_real_loss": [], "G_loss": []}
29
30    gp_weight = 10
31
32    try:
33        for iteration in range(1, opts.train_iters + 1):
34
35            # Reset data_iter for each epoch
36            if iteration % iter_per_epoch == 0:
37                train_iter = iter(dataloader)
38
39            real_images, real_labels = train_iter.next()
40            real_images, real_labels = to_var(real_images), to_var(real_labels).long().
squeeze()
41
42            # ones = Variable(torch.Tensor(real_images.shape[0]).float().cuda().fill_(1.0),
requires_grad=False)
43
44            for d_i in range(opts.d_train_iters):
45                d_optimizer.zero_grad()
46

```

```

47     # FILL THIS IN
48     # 1. Compute the discriminator loss on real images
49     D_real_loss = torch.mean((D(real_images) - 1) ** 2) / 2
50
51     # 2. Sample noise
52     noise = torch.randn_like(fixed_noise)
53
54     # 3. Generate fake images from the noise
55     fake_images = G(noise)
56
57     # 4. Compute the discriminator loss on the fake images
58     D_fake_loss = torch.mean(D(fake_images) ** 2) / 2
59
60     # ---- Gradient Penalty ----
61     if opts.gradient_penalty:
62         alpha = torch.rand(real_images.shape[0], 1, 1, 1)
63         alpha = alpha.expand_as(real_images).cuda()
64         interp_images = Variable(alpha * real_images.data + alpha * fake_images.
data, requires_grad=True).cuda()
65         D_interp_output = D(interp_images)
66
67         gradients = torch.autograd.grad(outputs=D_interp_output, inputs=
interp_images,
68                                         grad_outputs=torch.ones(D_interp_output.
size()).cuda(),
69                                         create_graph=True, retain_graph=True)[0]
70         gradients = gradients.view(real_images.shape[0], -1)
71         gradients_norm = torch.sqrt(torch.sum(gradients ** 2, dim=1) + 1e-12)
72
73         gp = gp_weight * gradients_norm.mean()
74     else:
75         gp = 0.0
76
77     # -----
78     # 5. Compute the total discriminator loss
79     D_total_loss = D_real_loss + D_fake_loss
80
81     D_total_loss.backward()
82     d_optimizer.step()
83
84     #####
85     ##          TRAIN THE GENERATOR          ##
86     #####
87
88     g_optimizer.zero_grad()
89
90     # FILL THIS IN
91     # 1. Sample noise
92     noise = torch.randn_like(fixed_noise)
93
94     # 2. Generate fake images from the noise
95     fake_images = G(noise)
96
97     # 3. Compute the generator loss
98     G_loss = torch.mean((D(fake_images) - 1)**2)
99
100    G_loss.backward()
101    g_optimizer.step()
102

```

```

103         # Print the log info
104         if iteration % opts.log_step == 0:
105             losses['iteration'].append(iteration)
106             losses['D_real_loss'].append(D_real_loss.item())
107             losses['D_fake_loss'].append(D_fake_loss.item())
108             losses['G_loss'].append(G_loss.item())
109             print('Iteration [{:4d}/{:4d}] | D_real_loss: {:.64f} | D_fake_loss: {:.64f}
110                   | G_loss: {:.64f}'.format(
111                         iteration, total_train_iters, D_real_loss.item(), D_fake_loss.item(),
112                         G_loss.item()))
113
114         # Save the generated samples
115         if iteration % opts.sample_every == 0:
116             gan_save_samples(G, fixed_noise, iteration, opts)
117
118         # Save the model parameters
119         if iteration % opts.checkpoint_every == 0:
120             gan_checkpoint(iteration, G, D, opts)
121
122     except KeyboardInterrupt:
123         print('Exiting early from training.')
124         return G, D
125
126     plt.figure()
127     plt.plot(losses['iteration'], losses['D_real_loss'], label='D_real')
128     plt.plot(losses['iteration'], losses['D_fake_loss'], label='D_fake')
129     plt.plot(losses['iteration'], losses['G_loss'], label='G')
130     plt.legend()
131     plt.savefig(os.path.join(opts.sample_dir, 'losses.png'))
132     plt.close()
133     return G, D

```

## 1.3 Experiments

# 2 Part 3: BigGAN [2pt]

## 2.1 BigGAN Experiments


### 2.1.1 Question 1

**Method** Given two classes with embeddings  $\mathbf{r}_{class1}$  and  $\mathbf{r}_{class2}$  in T-SNE space, to decide whether they are good candidate for linear interpolation, I looked at convex combinations of two classes' embeddings in figure 2.1. Let  $\Phi$  denote the set of convex combinations:

$$\Phi = \{\alpha \mathbf{r}_{class1} + (1 - \alpha) \mathbf{r}_{class2} : \alpha \in [0, 1]\} \quad (2.1)$$

Linear interpolations between these two classes are basically taking points in set  $\Phi$  and visualizing them. If there are many other classes' embeddings on or near the set  $\Phi$ , points in  $\Phi$  are likely to be associated with meaningful visualizations. Therefore, linear interpolations between these two classes are likely to generate meaningful results. Otherwise, these two classes are not good candidates for linear interpolation.

Figure 2.1: T-SNE Plot



**Good Candidate Match 1** There are many classes near the segment between **folding chair** and **vending machine** (right-up corner), so linear interpolation should be meaningful.

**Good Candidate Match 2** Similarly, there are many classes on the line between **bookshop** and **hot pot** (right-up corner), so they are ideal for linear interpolation.

**Bad Candidate Match 1** There are no other classes' embeddings are in between embeddings of **Go-Kart** and **Chessburger** in T-SNE plot, so they are not good match.

**Bad Candidate Pair 2** Similarly, **Ox** and **Chessburger** are not good match.

## 2.1.2 Question 2

### Implementation

```

1 #Linear interpolation between two class embeddings
2 def generate_linear_interpolate_sample(G, batch_size, class_label1, class_label2, alpha):
3     G.eval()
4     G.to(DEVICE)
5     with torch.no_grad():
6         z = torch.randn(batch_size, G.dim_z).to(DEVICE)
7         class1_emb = G.shared(torch.tensor(class_label1).to(DEVICE)*torch.ones((batch_size,))
8         ).to(DEVICE).long())
9         class2_emb = G.shared(torch.tensor(class_label2).to(DEVICE)*torch.ones((batch_size,))
10        ).to(DEVICE).long())
11
12     #####

```

```

11     ## FILL THIS IN: CREATE NEW EMBEDDING ##
12     #####
13     new_emb = alpha * class1_emb + (1 - alpha) * class2_emb
14
15     images = G(z, new_emb)
16     return images
17

```

## Linear Interpolation Results

Figure 2.2: Good Match 1: Folding Chair and Vending Machine



Figure 2.3: Good Match 2 Hot Pot and Bookshop



Figure 2.4: Bad Match 1 Go-Kart and Chessburger



Figure 2.5: Bad Match 2 Ox and Chessburger

